

# Open source monitoring and observability stack on Kubernetes

## Student Team

Pooja Srinivasan (1BM18CS069)  
Anusree Manoj K (1BM18CS017)  
Niha (1BM18CS060)  
Shikha N (1BM18CS149)

## Faculty Mentor

Dr. Nandini Vineeth  
Assistant Professor

## HPE Mentors

Divakar Padiyar  
Sonu Sudhakaran



kubernetes

# Agenda

- Introduction
- Abstract
- What is observability
- Pillars of observability
- Observability vs Monitoring
- Observability Use cases
- Observability solution with Prometheus, Grafana, Loki and Jaeger
- Demo
- Learnings
- Next steps

# Introduction

- **Aim** - Open source monitoring and observability stack on kubernetes.
- **Kubernetes** - open-source container orchestration platform.
- Designed to automate the **deployment, scaling, and management** of containerized applications.
- **Three pillars of observability - logs, metrics, and traces**  
When combined, they provide sufficient insights to monitor software at any scale.
- The project covers Monitoring, Alerting/Visualization, Log Aggregation/analytics, and Distributed systems tracing infrastructure which collectively make up observability.

# Abstract

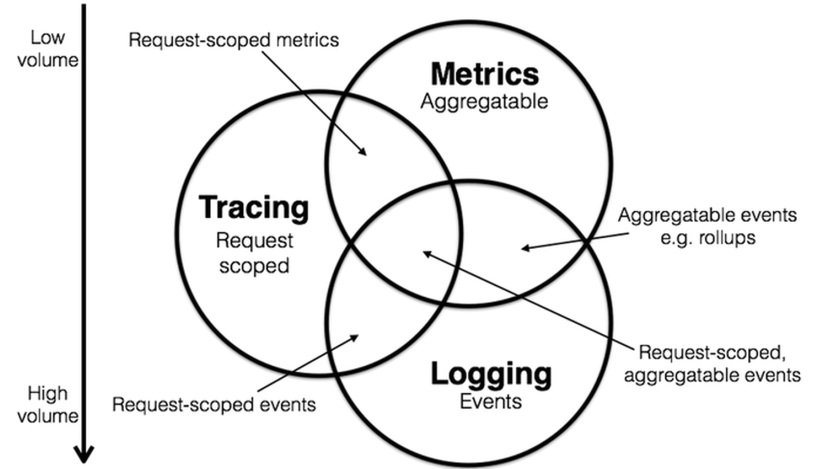
- Physical Server vs Virtual Machine vs Containers.
- Containerization and role of Kubernetes.
- Observability and Monitoring with respect to K8s.
- Setting up kubernetes tools - Micro K8s, Minikube, K3s.
- Using Grafana and Prometheus stack on the cluster to observe metrics of cluster.
- Configuring Alertmanager to receive alerts on slack channel.
- Logging and Tracing using Loki and Jaeger.

# What is Observability?

- A way to get insights into the whole infrastructure.
- Can explain any questions about what is happening on the inside of the system just by observing the outside of the system.
- Helps developers understand multi-layered architectures: what's slow, what's broken, and what needs to be done to improve performance.
- Creates an insight through an actionable knowledge of the whole environment by assembling all fragments from logs, monitoring tools and organizing them.

# Three Pillars Of Observability

- **Metrics:**  
These are numeric representation of data measured over intervals of time.
- **Logging:**  
They are discrete events and data in a structured textual form.
- **Tracing:**  
Represents consecutive events which reflect an end-to-end request path in a distributed system.



# Observability vs Monitoring

Observability	Monitoring
What is the system doing?	Is the system working?
Tells us <b>why</b> something goes wrong	Tells us <b>when</b> something went wrong
<b>Proactive</b> in nature	<b>Reactive</b> in nature
<b>Reduces the duration and impact</b> of incidents	Enables <b>quick response</b> when an incident occurs
Gain understanding <b>actively</b>	Consume information <b>passively</b>
Build to tame dynamic environments with changing complexity	Built to maintain static environments with little variation
Preferred by developers of systems with variability and <b>unknown permutations</b>	Used by developers of systems with little change and <b>known permutation</b>

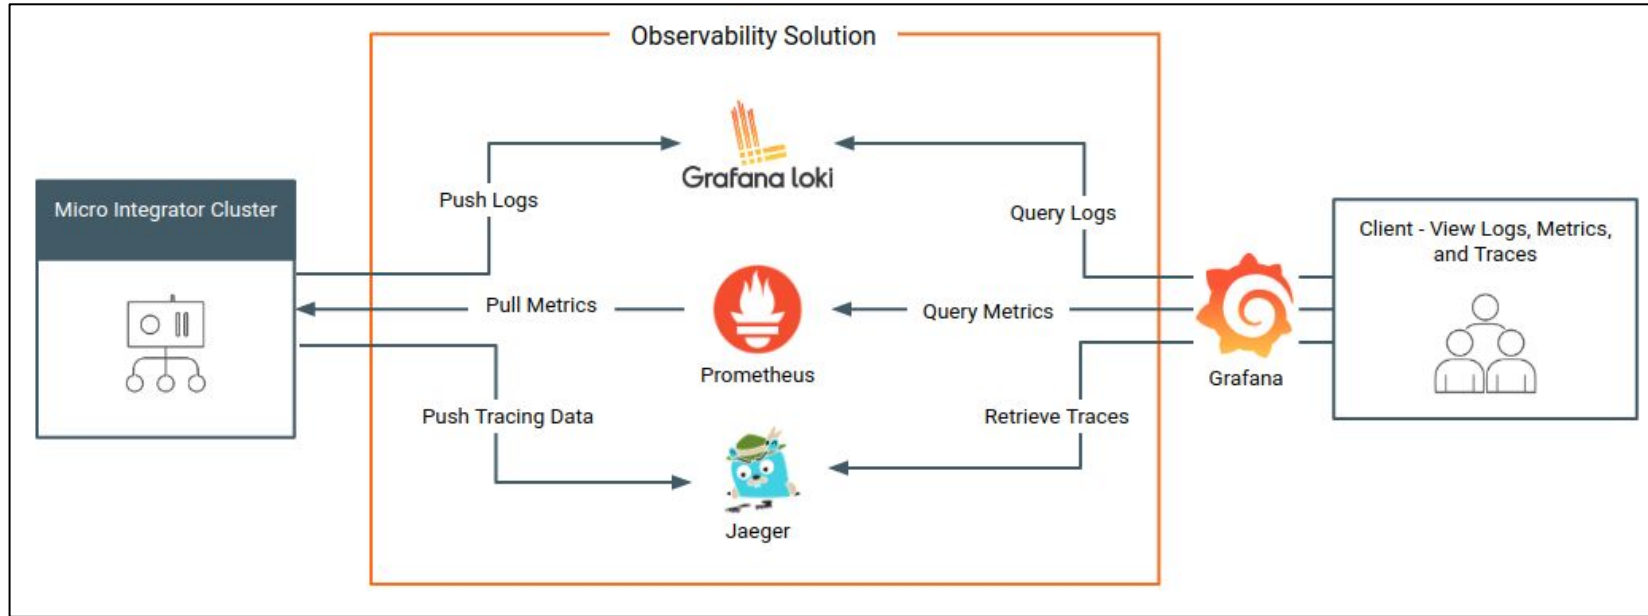


# Observability Use cases

- As an Infrastructure Admin, I would like the "Platform" services are monitored constantly for application runtime errors, re-curing operational / API / UI failures and generate alerts so that assigned / scheduled Support team gets notified immediately and is able to troubleshoot the error through Operations Console.
- As an Infrastructure Admin, I would like to configure Monitoring in the "Platform" to send Monitoring alerts to Operations console for any persistent failures in the System for timely response from the Operations team.
- As an SRE, I would like to monitor a gradual but consistent degradation of the Platform services performance / response times so that I can rectify any hardware resource related or scaling issues with the application and prevent missing any SLA(s).
- As an IT Operator, I would like to monitor cluster kube state and node metrics of Platform.
- As an IT Operator, I would like to monitor all the namespaces of the Platform.



# Observability solution with Prometheus, Grafana, Loki and Jaeger

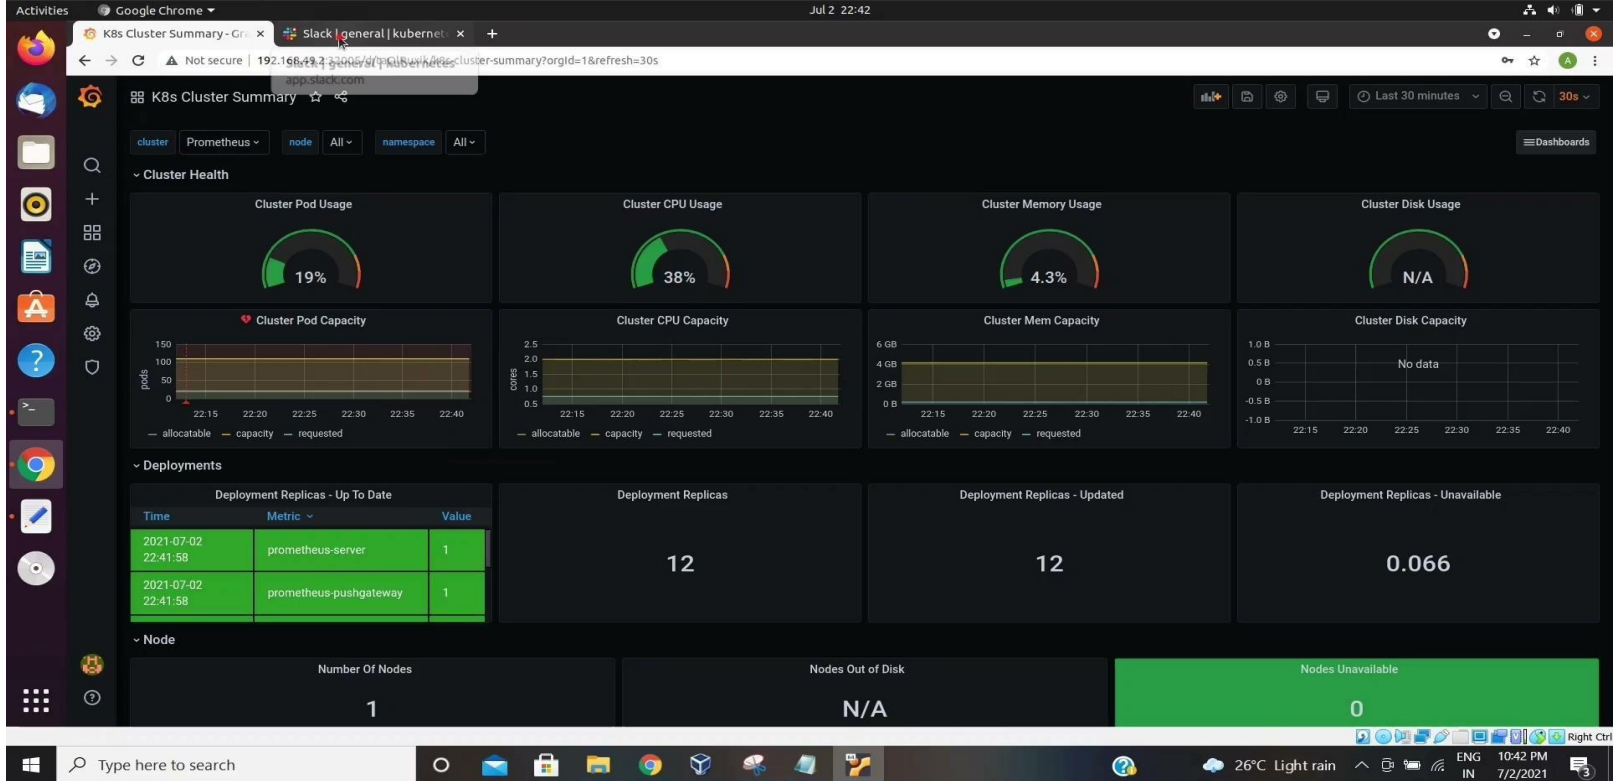


# Demo

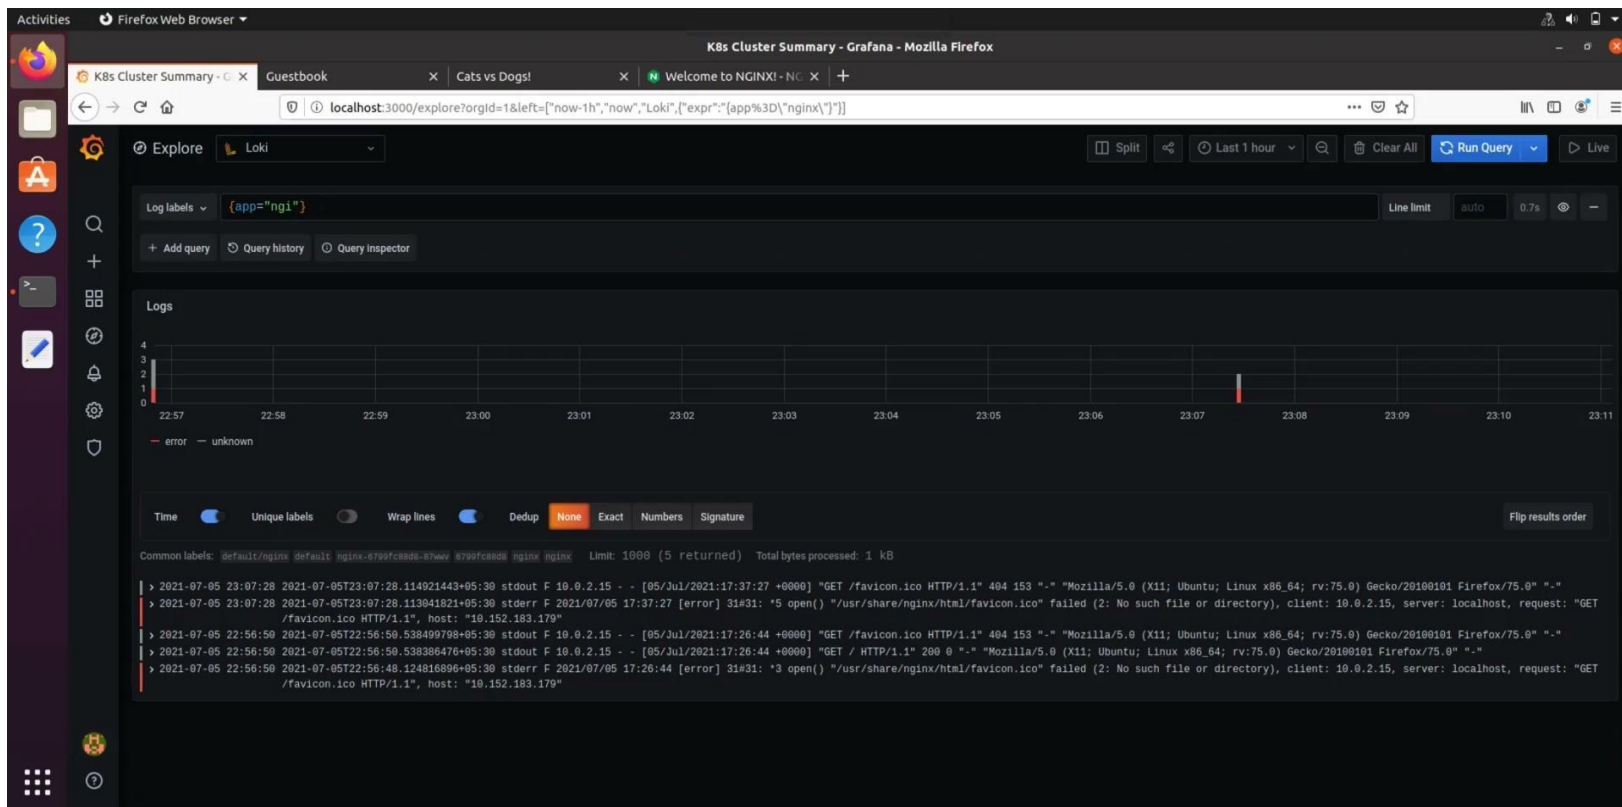


kubernetes [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help



# Demo



# Demo



kube1 [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Firefox Web Browser Jul 7 22:14

Kubernetes Dashboard HotROD - Rides On Demand Jaeger UI Home - Grafana

127.0.0.1:8001/api/v1/namespaces/kubernetes-dashboard/services/https:kubernetes-dashboard/proxy/#/pod?namespace=default

kubernetes default Search

Workloads > Pods

**Workloads**

- Cron Jobs
- Daemon Sets
- Deployments
- Jobs
- Pods**
- Replica Sets
- Replication Controllers
- Stateful Sets

**Service**

- Ingresses
- Services

**Config and Storage**

- Config Maps
- Persistent Volume Claims
- Secrets
- Storage Classes

**Cluster**

- Cluster Role Bindings
- Cluster Roles

**CPU Usage**

**Memory Usage**

**Pods**

Name	Namespace	Images	Labels	Node	Status	Restarts	CPU Usage (cores)	Memory Usage (bytes)	Created
jaeger-7dd885bd4b-kppv2	default	jaegertracing/all-in-one:latest	app: jaeger pod-template-hash: 7dd885bd4b	kube1	Running	7	0.00m	12.8Mi	7 days ago
hotrod-64d95fc446-xjds9	default	jaegertracing/example-hotrod:latest	app: hotrod pod-template-hash: 64d95fc446	kube1	Running	7	0.00m	7.3Mi	7 days ago
loki-0	default	grafana/loki:2.2.0	app: loki controller-revision-hash: loki-59dfd55c0b name: loki Show all	kube1	Running	16	0.00m	24.5Mi	7 days ago
			app: promtail						



# Learnings


- Importance of containerization and container orchestration.
- Importance of using Kubernetes.
- Using different tools to set up kubernetes clusters.
- Importance of observability and monitoring.
- Exploring different tools for observability and monitoring.
- Deploying multiple applications on a kubernetes cluster.
- Practical implementation of three pillars of observability i.e metrics, logging and tracing.
- Exploring multi-cluster observability.

# Next steps

- Creating a central dashboard for visualizing metrics of multiple clusters.
- Using Spark to process Big Data on Kubernetes.
- Implementation of load balancers using traefik.
- Monitoring and observability using EFK Stack.
- Observability using Open Telemetry.
- Setting up a High availability cluster using K3d.
- AI/ML stack on Kubernetes.

# CHALLENGES FACED

- Working remotely
- Insufficient hardware
- Ramp up time in working with various tools



# Any Questions?





# Thank You!